

# Modelos paramétricos sencillos para clasificación y regresión


MIAX-12, mayo 2024

# Contenidos

- Regresión lineal
- Regresión logística
- Problemas no lineales
- Regularización

# Regresión lineal


- **Problema:**  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_i \equiv$  vector de atributos
  - $t_i \equiv$  variable objetivo (target),  $t_i \in \mathbf{R}$
  - $N \equiv$  número de ejemplos/patrones
- **Objetivo:** predecir  $t$  a partir de  $\mathbf{x}$
- **Modelo:**  $y = f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
- **Función de coste:**  $L = \sum_i (y_i - t_i)^2$  (error cuadrático)



El modelo se entrena buscando el conjunto de parámetros  $\mathbf{w}$ ,  $b$  que minimiza la función de coste sobre los datos de entrenamiento.

# Regresión lineal 1D

x es un escalar



$$y = f(x) = wx + b$$

$$w = \frac{\text{cov}(x, t)}{\text{Var}(x)}$$
$$b = \bar{t} - w\bar{x}$$

# Regresión lineal multidimensional

$$\mathbf{x} = (x_1, x_2, \dots, x_d)$$

$$y = b + w_1 x_1 + w_2 x_2 + \dots + w_d x_d$$

Transformación

$$\left. \begin{aligned} \vec{x} &= (1, x_1, x_2, \dots, x_d) \\ y &= \sum_{j=0}^d w_j x_j \\ \vec{w} &= (w_0, w_1, \dots, w_d) \end{aligned} \right\} y = \vec{x} \cdot \vec{w}^t$$

# Regresión lineal multidimensional

$$\mathbf{x} = (x_1, x_2, \dots, x_d)$$

$$y = b + w_1 x_1 + w_2 x_2 + \dots w_d x_d$$

Matricialmente:

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1d} \\ 1 & x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & \dots & x_{Nd} \end{pmatrix}$$

$$\mathbf{Y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix}$$

$$\mathbf{T} = \begin{pmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{pmatrix}$$

Modelo:  $\mathbf{Y} = \mathbf{X}\mathbf{w}^T$

Solución:  $\mathbf{w}^T = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{T}$

# Ejemplos

## `sklearn.linear_model.LinearRegression`

```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True, copy_X=True, n_jobs=None, positive=False) \[source\]
```



Ordinary least squares Linear Regression.

`LinearRegression` fits a linear model with coefficients  $w = (w_1, \dots, w_p)$  to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LinearRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html)

Ver notebook [10\\_1\\_modelos\\_lineales.ipynb](#)

# Problemas no lineales




¿Podemos mantener el enfoque lineal?




# Problemas no lineales

1. Realizar una transformación no lineal de los atributos del problema



2. Resolver el problema linealmente en el nuevo espacio de atributos




# Problemas no lineales




$$\Phi(x) = (x, x^2)$$



# Problemas no lineales




$$\Phi(x) = (x, x^2, x^3, x^4)$$



# Sobre la complejidad del modelo

- ¿Qué complejidad elegir?
- ¿Deberíamos elegir un modelo lo más complejo posible? (Por si acaso...)



$$\Phi(x) = (x)$$


$$\Phi(x) = (x, x^2, x^3, x^4)$$

$$\Phi(x) = (x, x^2, x^3, x^4, x^5, x^6, x^7)$$


}

?


# Sobre la complejidad del modelo



$$\Phi(x) = (x)$$



$$\Phi(x) = (x, x^2, x^3, x^4, x^5, x^6, x^7)$$



# El dilema sesgo-varianza

Fuentes de error en un modelo:

- **Sesgo**: relacionado con la capacidad del modelo para ajustarse a los datos
- **Varianza**: relacionado con la sensibilidad del modelo a las variaciones en los datos
- **Error intrínseco**: inherente a los datos

$$\underbrace{E_{\mathbf{x},y,D} \left[ (h_D(\mathbf{x}) - y)^2 \right]}_{\text{Expected Test Error}} = \underbrace{E_{\mathbf{x},D} \left[ (h_D(\mathbf{x}) - \bar{h}(\mathbf{x}))^2 \right]}_{\text{Variance}} + \underbrace{E_{\mathbf{x},y} \left[ (\bar{y}(\mathbf{x}) - y)^2 \right]}_{\text{Noise}} + \underbrace{E_{\mathbf{x}} \left[ (\bar{h}(\mathbf{x}) - \bar{y}(\mathbf{x}))^2 \right]}_{\text{Bias}^2}$$

Normalmente los modelos más complejos reducen el sesgo, pero aumentan la varianza → Mayor riesgo de **overfitting**

# El dilema sesgo-varianza



**Variance:** Captures how much your classifier changes if you train on a different training set. How "over-specialized" is your classifier to a particular training set (overfitting)? If we have the best possible model for our training data, how far off are we from the average classifier?

**Bias:** What is the inherent error that you obtain from your classifier even with infinite training data? This is due to your classifier being "biased" to a particular kind of solution (e.g. linear classifier). In other words, bias is inherent to your model.

**Noise:** How big is the data-intrinsic noise? This error measures ambiguity due to your data distribution and feature representation. You can never beat this, it is an aspect of the data.

<https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote12.html>

# El dilema sesgo-varianza



<http://scott.fortmann-roe.com/docs/BiasVariance.html>



# Regularización

Incluir en la función de coste términos que **penalizan la complejidad** del modelo

$$J = \underbrace{C(y, t)}_{\text{error}} + \lambda \cdot \underbrace{\text{COMPLEJIDAD}}_{\|\bar{w}\|^2}$$

¿Cómo medir la complejidad de un modelo?

# Regularización L2 (Ridge)

## `sklearn.linear_model.Ridge`

```
class sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, copy_X=True, max_iter=None, tol=0.0001, solver='auto',  
positive=False, random_state=None)
```

[\[source\]](#)

Linear least squares with l2 regularization.

Minimizes the objective function:

$$||y - Xw||^2_2 + \alpha * ||w||^2_2$$

This model solves a regression model where the loss function is the linear least squares function and regularization is given by the l2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression (i.e., when  $y$  is a 2d-array of shape  $(n\_samples, n\_targets)$ ).

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.Ridge.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html)

# Regularización L1 (Lasso)

## `sklearn.linear_model.Lasso`

```
class sklearn.linear_model.Lasso(alpha=1.0, *, fit_intercept=True, precompute=False, copy_X=True, max_iter=1000, tol=0.0001, warm_start=False, positive=False, random_state=None, selection='cyclic')
```

[\[source\]](#)

Linear Model trained with L1 prior as regularizer (aka the Lasso).

The optimization objective for Lasso is:

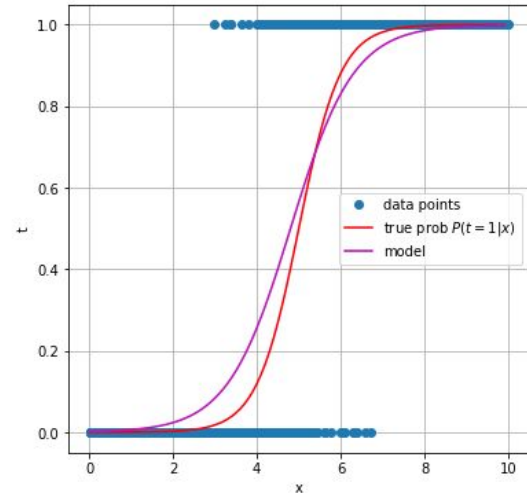
$$(1 / (2 * n\_samples)) * ||y - Xw||^2 + alpha * ||w||_1$$

Technically the Lasso model is optimizing the same objective function as the Elastic Net with `l1_ratio=1.0` (no L2 penalty).

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.Lasso.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Lasso.html)


# Regresión logística (clasificación)

- **Problema:**  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_i \equiv$  vector de atributos
  - $t_i \equiv$  variable objetivo (target),  $t_i \in \{0, 1\}$
  - $N \equiv$  número de ejemplos/patrones
- **Objetivo:** predecir  $t$  a partir de  $\mathbf{x}$
- **Modelo:**  $y = f(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + b)$
- **Función de coste:**  $L = -\sum_i [t_i \log(y_i) + (1-t_i) \log(1-y_i)]$  (cross-entropy)



El modelo se entrena buscando el conjunto de parámetros  $\mathbf{w}$ ,  $b$  que minimizan la función de coste sobre los datos de entrenamiento.

# Función logística (sigmoide)



- Siempre toma valores entre 0 y 1
- La interpretamos como la probabilidad que da el modelo a la clase 1

# Regresión logística en scikit learn

## `sklearn.linear_model.LogisticRegression`

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

[\[source\]](#)

Logistic Regression (aka logit, MaxEnt) classifier.

In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme if the 'multi\_class' option is set to 'ovr', and uses the cross-entropy loss if the 'multi\_class' option is set to 'multinomial'. (Currently the 'multinomial' option is supported only by the 'lbfgs', 'sag', 'saga' and 'newton-cg' solvers.)

This class implements regularized logistic regression using the 'liblinear' library, 'newton-cg', 'sag', 'saga' and 'lbfgs' solvers. **Note that regularization is applied by default.** It can handle both dense and sparse input. Use C-ordered arrays or CSR matrices containing 64-bit floats for optimal performance; any other input format will be converted (and copied).

The 'newton-cg', 'sag', and 'lbfgs' solvers support only L2 regularization with primal formulation, or no regularization. The 'liblinear' solver supports both L1 and L2 regularization, with a dual formulation only for the L2 penalty. The Elastic-Net regularization is only supported by the 'saga' solver.

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

# Ejemplos

Ver notebook [\*10\\_1\\_modelos\\_lineales.ipynb\*](#)

# Siguientes pasos

- Métodos de kernel
- Regresión lineal basada en kernels (kernel ridge)
- Support vector machines